

Introduction to snakemake

A workflow management system for bioinformatics

Verena Kutschera (verena.kutschera@scilifelab.se)

National Bioinformatics Infrastructure Sweden (NBIS)

Science for Life Laboratory Solna & Dept. for Biochemistry and Biophysics, Stockholm university

2019-09-03



nextflow



Luigi



bpipe



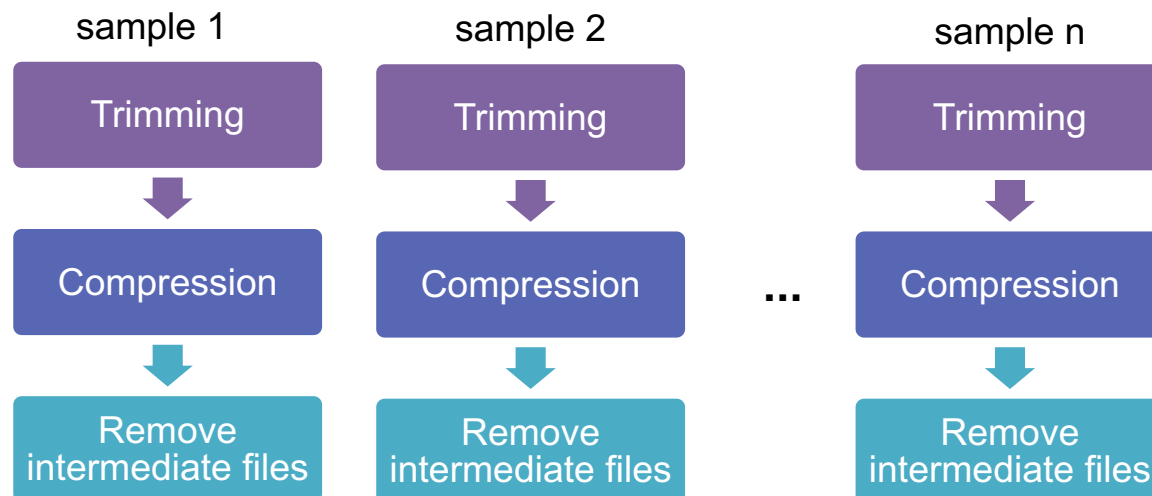
Ruffus

- Bash code

```
for sample in $(ls *.fastq | sed 's/\.fastq//')
do
  # Trim fastq file
  trim_galore --illumina ${sample}.fastq > ${id}.trimmed.fastq

  # Compress fastq file
  gzip -c ${sample}.trimmed.fastq > ${sample}.trimmed.fastq.gz


  # Remove intermediate files
  rm ${sample}.trimmed.fastq
done
```



- What kinds of problems have you encountered when running your analyses using bash scripts similar to the example?

- Interruption → continue where left off
- New samples → do not rerun everything
- Updated dependency → update downstream files, too

“There is a learning curve to adopting workflow languages. But, [...] “the energy that you expend learning is more than made up for by the energy you save in having your code be reproducible.” (Brian Naughton)”



<https://www.nature.com/articles/d41586-019-02619-z>



Python



GNU Make



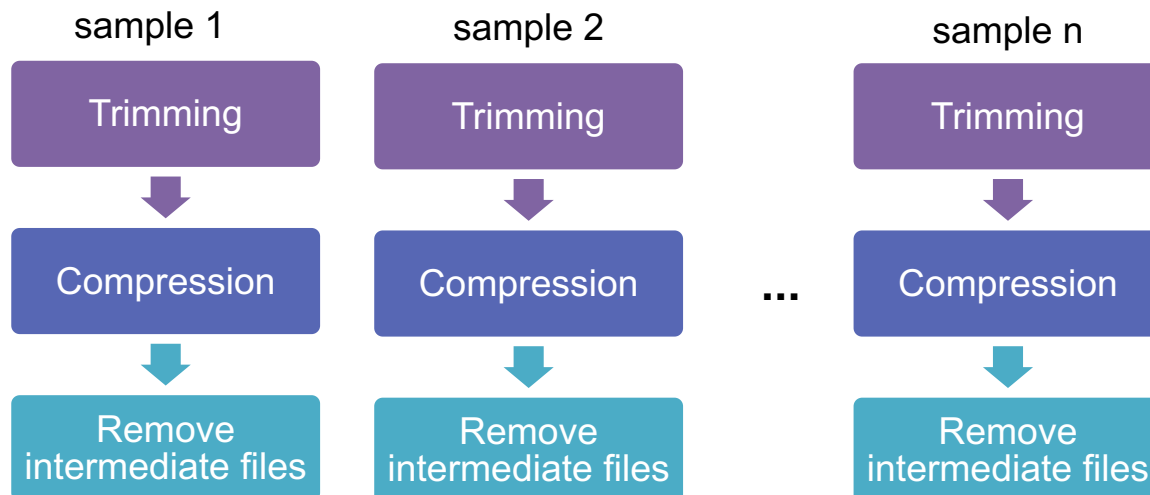
Snakemake

- Bash code

```
for sample in $(ls *.fastq | sed 's/\.fastq//')
do
  # Trim fastq file
  trim_galore --illumina ${sample}.fastq > ${id}.trimmed.fastq

  # Compress fastq file
  gzip -c ${sample}.trimmed.fastq > ${sample}.trimmed.fastq.gz

  # Remove intermediate files
  rm ${sample}.trimmed.fastq
done
```



Snakemake

```
rule trim_fastq:
    input: "{prefix}.fastq"
    output: temp("{prefix}.trimmed.fastq")
    shell:
        "trim_galore --illumina {input} > {output}"

rule gzip:
    input: "{prefix}"
    output: "{prefix}.gz"
    shell:
        "gzip -c {input} > {output}"
```

```
$ snakemake {a,b}.trimmed.fastq.gz
Building DAG of jobs...
Using shell: /bin/bash
Provided cores: 1
Rules claiming more threads will be scaled down.
Job counts:
    count jobs
    2      gzip
    2      trim_fastq
    4
rule trim_fastq:
    input: b.fastq
    output: b.trimmed.fastq
    jobid: 2
    wildcards: prefix=b
Finished job 2.
1 of 4 steps (25%) done
rule gzip:
    input: b.trimmed.fastq
    output: b.trimmed.fastq.gz
    jobid: 0
    wildcards: prefix=b.trimmed.fastq
Removing temporary output file b.trimmed.fastq.
Finished job 0.
2 of 4 steps (50%) done
rule trim_fastq:
    input: a.fastq
    output: a.trimmed.fastq
    jobid: 3
    wildcards: prefix=a
Finished job 3.
3 of 4 steps (75%) done
rule gzip:
    input: a.trimmed.fastq
    output: a.trimmed.fastq.gz
    jobid: 1
    wildcards: prefix=a.trimmed.fastq
Removing temporary output file a.trimmed.fastq.
Finished job 1.
4 of 4 steps (100%) done
```

Snakemake

```
rule trim_fastq:
    input: "{prefix}.fastq"
    output: temp("{prefix}.trimmed.fastq")
    shell:
        "trim_galore --illumina {input} > {output}"

rule gzip:
    input: "{prefix}"
    output: "{prefix}.gz"
    shell:
        "gzip -c {input} > {output}"
```

What happens if we add
another sample?

```
$ snakemake {a,b}.trimmed.fastq.gz
Building DAG of jobs...
Using shell: /bin/bash
Provided cores: 1
Rules claiming more threads will be scaled down.
Job counts:
    count jobs
    2      gzip
    2      trim_fastq
    4
rule trim_fastq:
    input: b.fastq
    output: b.trimmed.fastq
    jobid: 2
    wildcards: prefix=b
Finished job 2.
1 of 4 steps (25%) done
rule gzip:
    input: b.trimmed.fastq
    output: b.trimmed.fastq.gz
    jobid: 0
    wildcards: prefix=b.trimmed.fastq
Removing temporary output file b.trimmed.fastq.
Finished job 0.
2 of 4 steps (50%) done
rule trim_fastq:
    input: a.fastq
    output: a.trimmed.fastq
    jobid: 3
    wildcards: prefix=a
Finished job 3.
3 of 4 steps (75%) done
rule gzip:
    input: a.trimmed.fastq
    output: a.trimmed.fastq.gz
    jobid: 1
    wildcards: prefix=a.trimmed.fastq
Removing temporary output file a.trimmed.fastq.
Finished job 1.
4 of 4 steps (100%) done
```

Snakemake

```
rule trim_fastq:
    input: "{prefix}.fastq"
    output: temp("{prefix}.trimmed.fastq")
    shell:
        "trim_galore --illumina {input} > {output}"

rule gzip:
    input: "{prefix}"
    output: "{prefix}.gz"
    shell:
        "gzip -c {input} > {output}"
```

```
$ snakemake {a,b,c}.trimmed.fastq.gz
Building DAG of jobs...
Using shell: /bin/bash
Provided cores: 1
Rules claiming more threads will be scaled down.
Job counts:
```

```
    count jobs
      1      gzip
      1      trim_fastq
      2
```

```
[Thu Aug 15 15:29:56 2019]
```

```
rule trim_fastq:
    input: c.fastq
    output: c.trimmed.fastq
    jobid: 5
    wildcards: prefix=c
```

```
[Thu Aug 15 15:29:56 2019]
```

```
Finished job 5.
1 of 2 steps (50%) done
```

```
[Thu Aug 15 15:29:56 2019]
```

```
rule gzip:
    input: c.trimmed.fastq
    output: c.trimmed.fastq.gz
    jobid: 2
    wildcards: prefix=c.trimmed.fastq
```

```
Removing temporary output file c.trimmed.fastq.
```

```
[Thu Aug 15 15:29:56 2019]
```

```
Finished job 2.
2 of 2 steps (100%) done
```

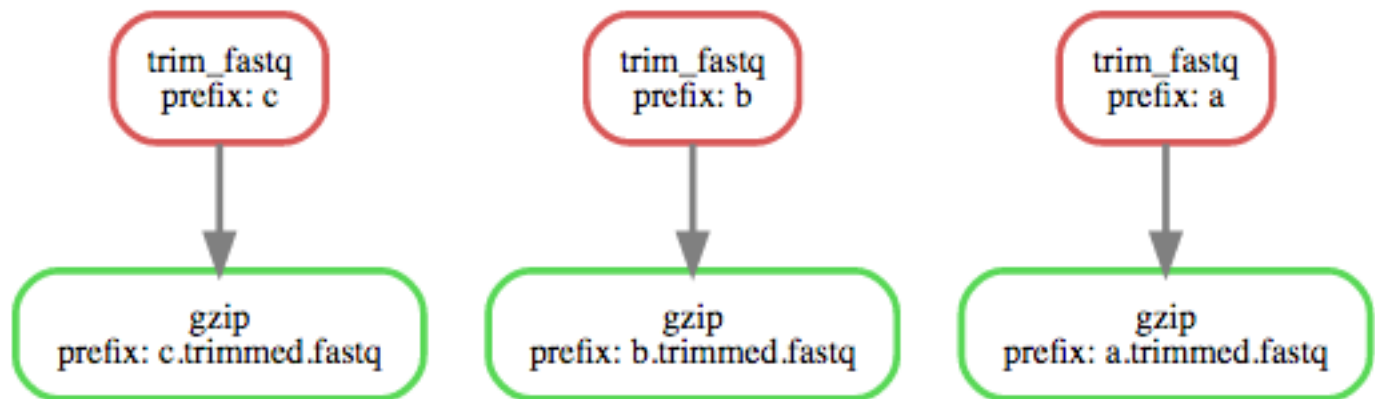
```
Complete log:
```

```
/Users/verenakutschera/nobackup/misc/teaching/2019
0903_snakemake_intro_EBC/workflows/.snakemake/log/
2019-08-15T152956.351207.snakemake.log
```

```
rule trim_fastq:
  input: "{prefix}.fastq"
  output: temp("{prefix}.trimmed.fastq")
  shell:
    "trim_galore --illumina {input} > {output}"

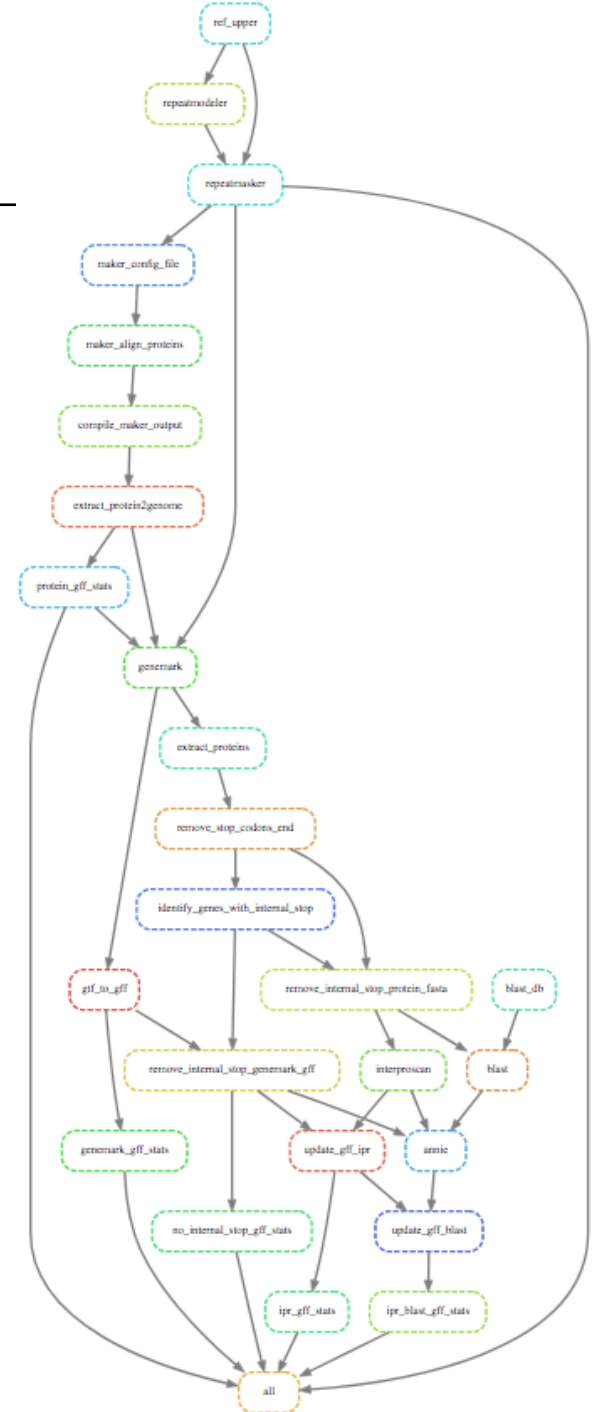
rule gzip:
  input: "{prefix}"
  output: "{prefix}.gz"
  shell:
    "gzip -c {input} > {output}"
```

- Snakemake keeps track of when files were generated and by which rules

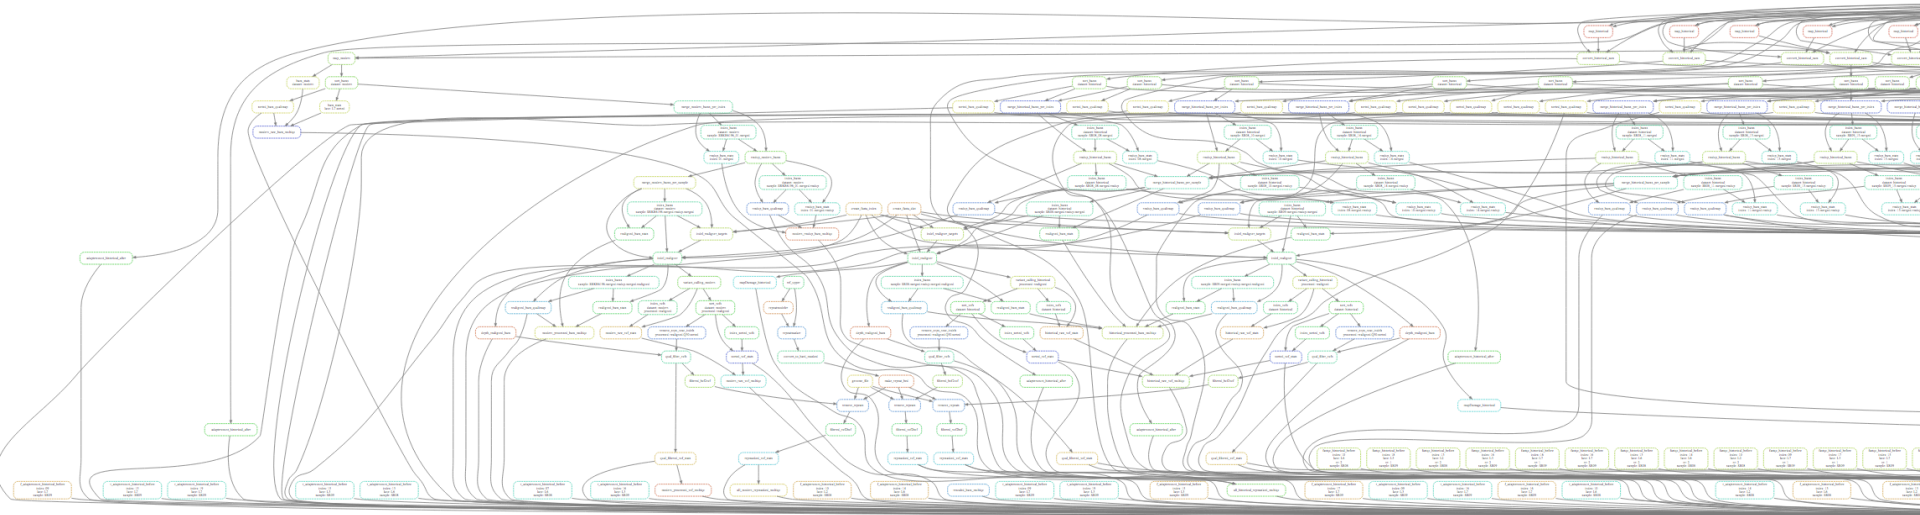


Rule graphs

- Workflows can get complex...



- Workflows can get complex...



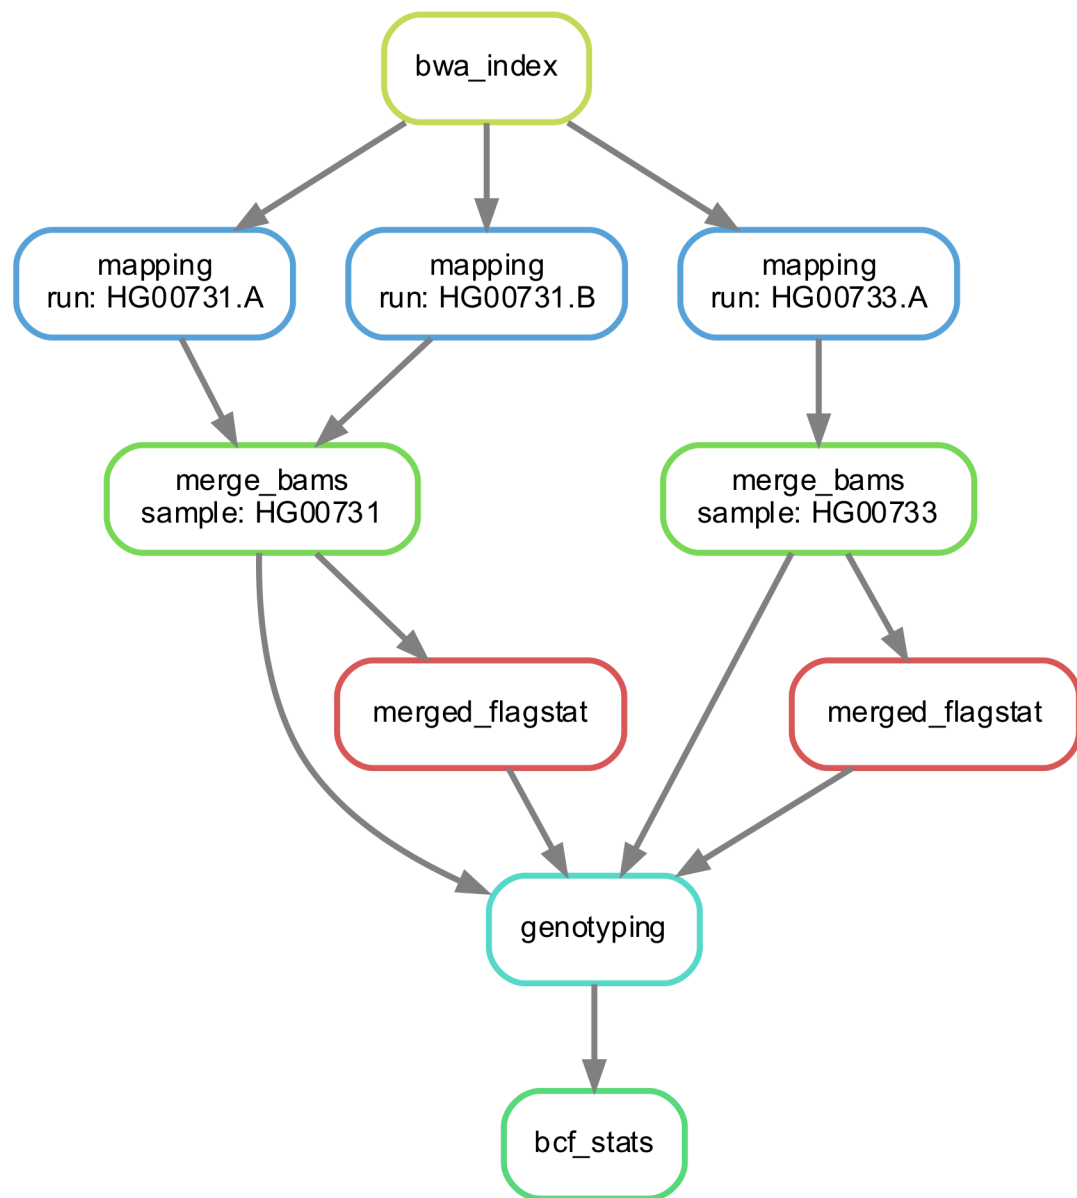
- Structured pipeline
- Ideal for parallel problems (i.e. analysis of many samples)
- Dependencies are handled
- Execution order is inferred from input and output file names
- Reentrance into analysis is possible
- New upstream files are automatically analyzed
- Logging system to follow the status

- Easy handling on a cluster
- Integration of various code, e.g. R, bash and python
- Rule syntax very flexible and adaptable
- Python
 - Glue language
 - Snakemake is written in python
 - Snakemake syntax is inspired by python
- Active development and large community

- Snakefiles are no python modules
- Errors sometimes cryptic and difficult to track down
- Risk to end up in an infinite loop or to delete the content of a file
 - Do a lot of testing
 - Backup the data
- Rule syntax is very flexible and adaptable
 - Different developers will come up with different solutions for the same problem – difficult to share workflows

- Convert bash script to snakemake workflow
- Run the workflow on UPPMAX
- Final workflow will be on data science GitHub repository (along with this presentation)

- FastQ files from two samples
 - HG00731 (run A and B)
 - HG00733 (run A)
- Map to reference genome
- Merge BAM files per sample
- Get BAM file statistics
- Joint genotyping
- Get VCF file statistics



- Snakefile
- Basic rule syntax
- Python functions in snakemake workflows
- Input lists
- “Expand” statement
- Dry mode
- Rule graphs (dag files)
- Target rule (“all”)
- Temporary files

- Tmux or screen to send snakemake to the background
 - tmux: check e.g. <https://danielmiessler.com/study/tmux/>
 - Write the standard output to a file with “&>” if you are running a long workflow or many samples at once

- Snakemake with cluster configuration sends each rule as job to the slurm system
 - cluster.json or cluster.yaml with slurm parameters for each rule
 - To use yaml on Rackham, install the python package PyYAML for python version 3.6.0 or higher
(<https://www.uppmax.uu.se/support/user-guides/python-modules-guide/>)
 - Start snakemake with this command to point to the cluster.yaml file:

```
$ snakemake -j 999 --cluster-config cluster.yaml \  
--cluster "sbatch -A {cluster.account} \  
-p {cluster.partition} -n {cluster.n} -t {cluster.time}"
```

- How do I cancel snakemake when it's running on the slurm system?
 1. cancel jobs (scancel)
 2. cancel snakemake process in tmux/screen session with CTRL+C

- How do I cancel snakemake when it's running on the slurm system?
 1. cancel jobs (scancel)
 2. cancel snakemake process in tmux/screen session with CTRL+C
- What do I do if a specific rule/job failed?
 - "Error in rule XYZ" in standard output
 - Look for the rule and the line "cluster_jobid: Submitted batch job 1234567" in standard output
 - Slurm output "slurm-1234567.out" will hopefully contain more info

- How do I cancel snakemake when it's running on the slurm system?
 1. cancel jobs (scancel)
 2. cancel snakemake process in tmux/screen session with CTRL+C
- What do I do if a specific rule/job failed?
 - "Error in rule XYZ" in standard output
 - look for the rule and the line "cluster_jobid: Submitted batch job 1234567" in standard output
 - slurm output "slurm-1234567.out" will hopefully contain more info
- The jobs submitted by snakemake are gone from the list of running jobs, but snakemake is still running. What happened, and how do I fix it?
 - snakemake and slurm are communicating about running jobs, except for jobs failing due to timeout
 - cancel the snakemake process with CTRL+C and restart workflow

- Threads
- Cluster configuration

- Snakemake sends many, small jobs to the slurm system
 - Define group for execution to send several rules in one job to slurm

```
samples = [1,2,3,4,5]

rule a:
    input:
        "a/{sample}.out"
    output:
        "a/{sample}.edits.out"
    group: "mygroup"
    shell:
        "touch {output}"

rule b:
    input:
        "a/{sample}.edits.out"
    output:
        "b/{sample}.out"
    group: "mygroup"
    shell:
        "touch {output}"

rule c:
    input:
        expand("b/{sample}.out", sample=samples)
    output:
        "test.out"
    shell:
        "touch {output}"
```

- Rules a and b are run as one job per sample in the group “mygroup”
- Rule c connects the other jobs
 - As part of “mygroup”, everything would be submitted as one group

- Snakemake sends many, small jobs to the slurm system
 - Define group for execution to send several rules in one job to slurm
 - Define local rules

```
localrules: all, foo  
  
rule all:  
    input: ...  
  
rule foo:  
    ...  
  
rule bar:  
    ...
```

- Rules all and foo will be executed on the login node
- Rule foo will be submitted as a job



- Defining groups for execution
- Local rules
- Config files
- Non-file parameters for rules
- Running python code instead of a shell command
- Log files

- Snakemake documentation on readthedocs:
 - <https://snakemake.readthedocs.io/en/stable/#>
- Stackoverflow:
 - <https://stackoverflow.com/questions/tagged/snakemake>
- NBIS bioinformatics drop-in:
 - <https://nbis.se/events/>
 - Uppsala: every Thursday 10-11h in Navet, BMC Uppsala E10:3

- Snakemake documentation:
 - <https://snakemake.readthedocs.io/en/stable/tutorial/tutorial.html>
- NBIS reproducible research workshop:
 - <https://nbis-reproducible-research.readthedocs.io/en/latest/snakemake/>

- NBIS reproducible research workshop:
 - <https://nbis-reproducible-research.readthedocs.io/en/latest/snakemake/>

Tools for reproducible research

Share with others:  

Course open for PhD students (prioritized), postdocs, researchers and other interested in making their computational analysis more reproducible.

Important dates

Application open: **August 28**

Application deadline: **October 22**

Confirmation to accepted participants: **October 28**

Date and time

2019-11-18 - 2019-11-20

Location

Chalmers University of Technology,
Kemivägen 10, SE-412 96,

- IN THE MAKING: 2-day NBIS snakemake workshop, fall 2020

- RNA-seq:
 - <https://github.com/crazyhottommy/RNA-seq-analysis/tree/master/RNA-seq-snakemake-pipeline>
 - <https://github.com/slowkow/snakefiles>
- GATK variant calling:
 - <https://github.com/snakemake-workflows/dna-seq-gatk-variant-calling>
- Samtools variant calling (incl. Docker / Singularity):
 - https://github.com/fbartusch/snakemake_tutorial
- ACCEL amplicon trimming:
 - <https://github.com/snakemake-workflows/accel-amplicon-trimming>
- Single-cell drop-seq:
 - <https://github.com/snakemake-workflows/single-cell-drop-seq>

- Presentations by NBIS colleagues Leif Wigge & Rasmus Ågren, Per Unneberg
- Snakemake demo inspired by material from Per Unneberg
- https://groupes.renater.fr/sympa/d_read/bios4biol/2017/VisioCati2017/slides_snakemake.pdf
- Snakemake tutorial & documentation:
<https://snakemake.readthedocs.io/en/stable/index.html>